

Formal Verification of a Flash Memory Device Driver - an Experience Report

Moonzoo Kim, Yunho Kim

Provable Software Lab. CS Dept. KAIST



Yunja Choi

School of EECS, Kyungpook National Univ.

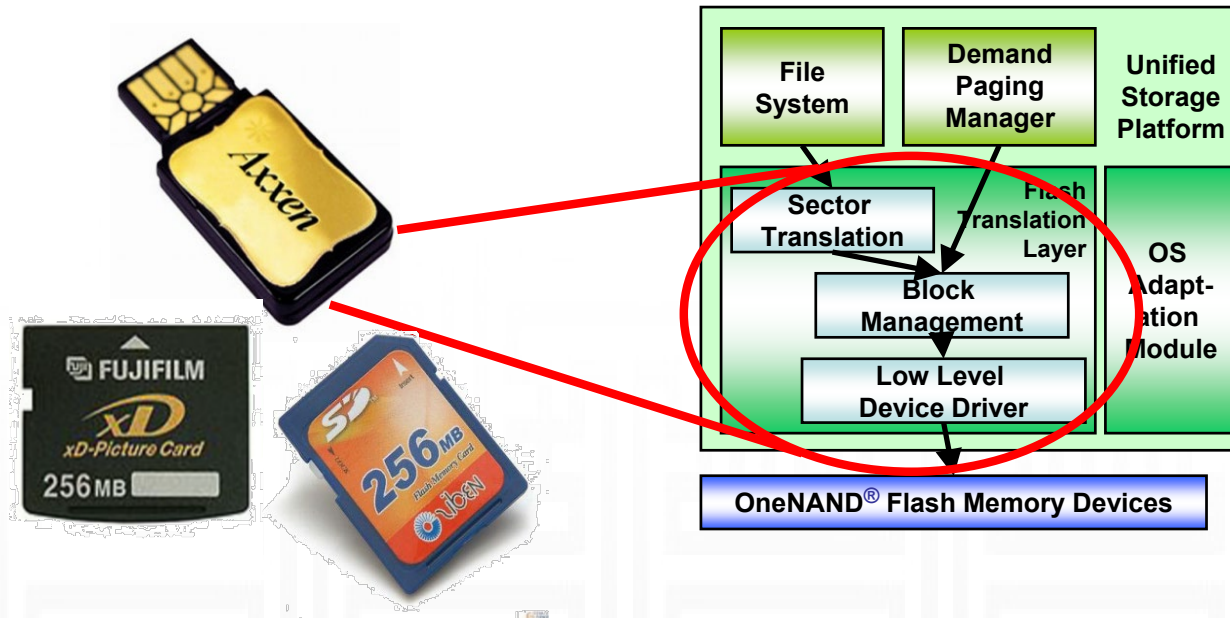


Hotae Kim

Samsung Electronics



Summary of the Talk



- In 2007, Samsung requested to debug the device driver for the Samsung OneNAND™ flash memory, by using model checkers, for 6 months. This presentation describes a part of the result from the project.

Overview

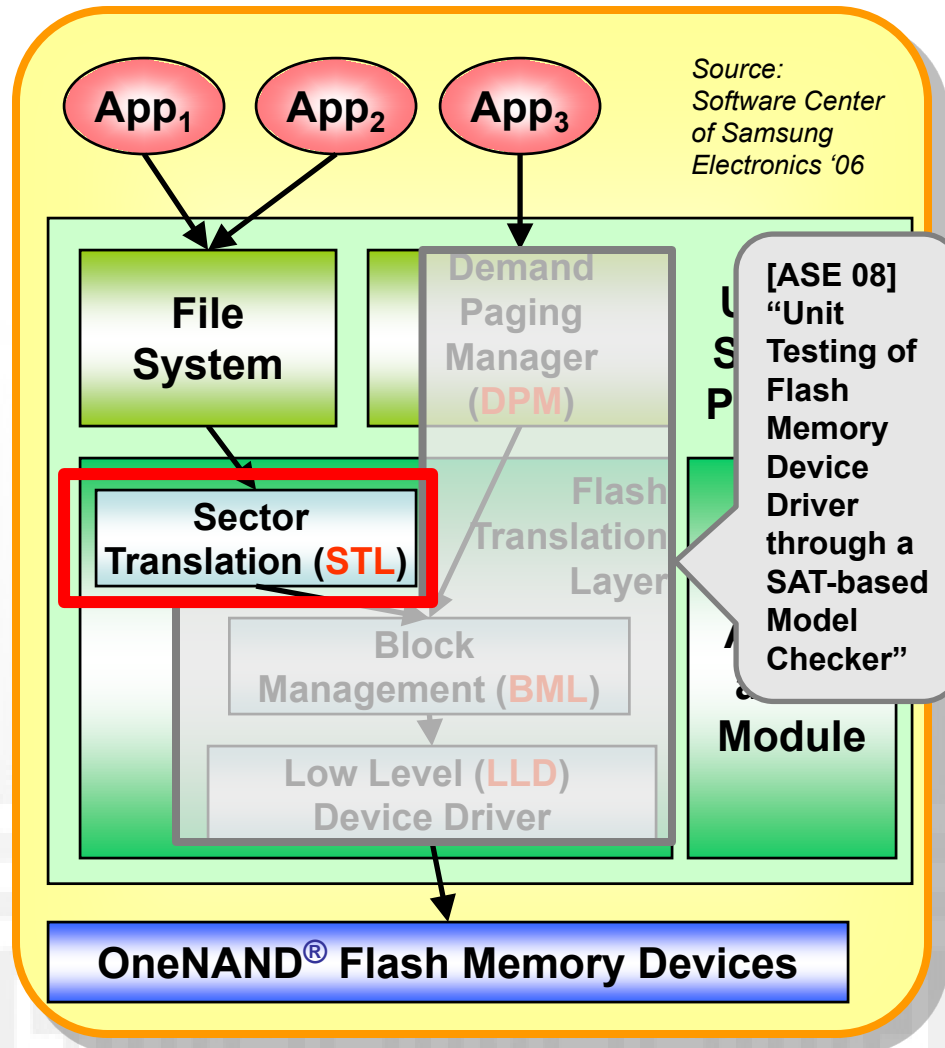
- **Background**
 - Overview of the Unified Storage Platform (USP)
 - Sector Translation Layer (STL)
 - Multi-Sector Read operation (MSR)
- **Model Checking MSR**
 - Reports on the following three aspects
 - Target system modeling
 - Environment modeling
 - Performance analysis on the verification
- **Three different types of model checkers are used**
 - BDD based symbolic model checking (NuSMV)
 - Explicit model checking (Spin)
 - C-bounded model checking (CBMC)

PART I: Background

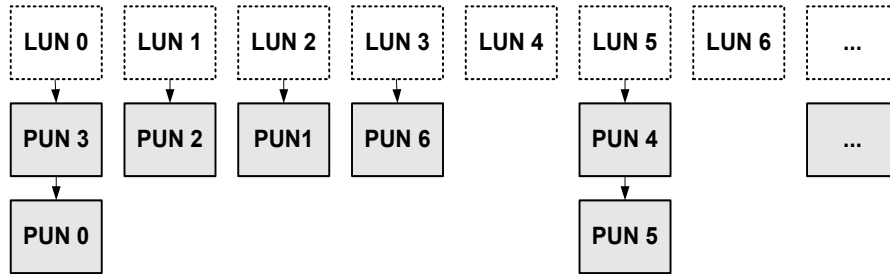
- **Unified Storage Platform (USP)**
 - Block diagram
 - Code statistics
- **Logical-to-physical sector translation**
 - Example of possible data distributions
- **Multi-Sector Read operation (MSR)**
 - Pseudo structure

Overview of the OneNAND[®] Flash Memory

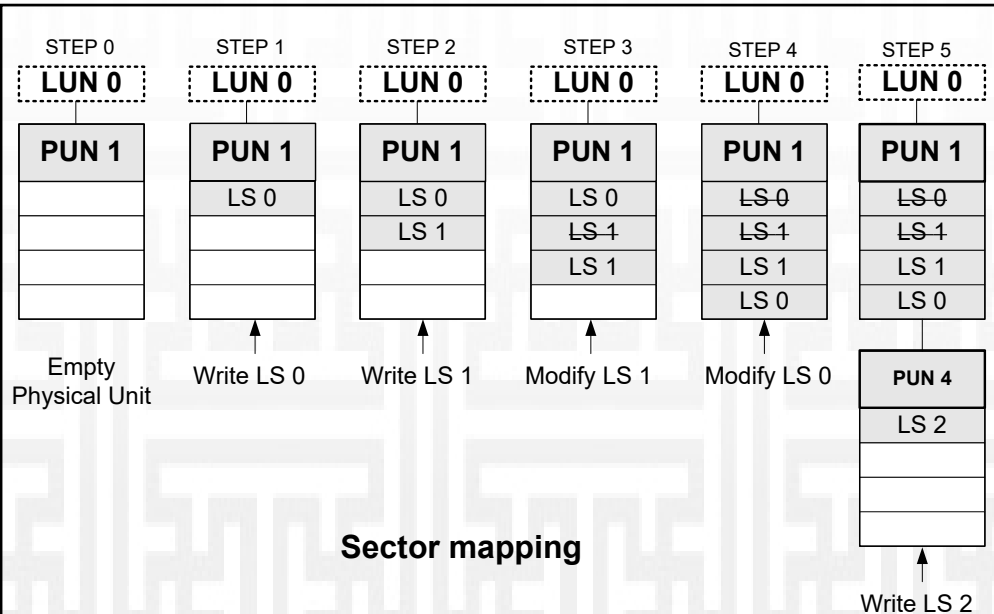
- **Characteristics of OneNAND[®]**
 - Each memory cell can be written limited number of times only
 - Logical-to-physical sector mapping
 - Bad block management
 - Wear-leveling
 - Performance enhancement
 - Multi-sector read/write
 - Asynchronous operations
 - Deferred operation result check



Logical to Physical Sector Mapping



1:N mapping from a LUN to PUNs



Sector mapping

SAM1

Logical offset	Physical offset
0	3
1	2
2	
3	

LUN 0

PUN 1
LS 0
LS 4
LS 1
LS 0

SAM4

Logical offset	Physical offset
0	
1	
2	0
3	

LUN 0

PUN 4
LS 2

Sector Allocation Map (SAM)

- In flash memory, logical data are distributed over physical sectors.

Examples of Possible Data Distribution

	SAM0~SAM4				PU0~PU4			
Sector 0	1			0			E	
Sector 1		1		1	A	B		F
Sector 2		2				C		
Sector 3			3				D	

(a) A distribution of
“ABCDEF”

	SAM0~SAM4				PU0~PU4			
		3		3	B			
0			2			D		
			3				F	
1					A	C		E

(b) Another distribution of
“ABCDEF”

	SAM0~SAM4				PU0~PU4			
	1			0			B	
		1		1	F	E		A
		2				D		
			3				C	

(c) A distribution of
“FEDCBA”

- **Assumptions**
 - there are 5 physical units
 - each unit has 4 sectors
 - each sector is 1 byte long

Multi-Sector Read Operations (MSR)

```
1032 nSamIdx = (UINT16)(nLsn % pstSHPC->nLogSctsPerUnit);
1033
1034 while (nNumOfScts > 0)
1035 {
1036     pstNew = pstSMC->pstLogUnitInfo[nLun].pstVirUnitInfo;
1037
1038     /* get the number of logical sectors to be read in a current logical uni
1039     nReadScts = ((pstSHPC->nLogSctsPerUnit - nSamIdx) > nNumOfScts) ? nNumO
1040     (pstSH
1041
1042     /* update nNumOfScts */
1043     nNumOfScts -= nReadScts;
1044
1045     if (pstNew != NULL)
1046     {
1047         /* construct SAM table */
1048         if (_ConstructSam(pstSMC, nLun, STL_LRU_POLICY) != STL_SUCCESS)
1049         {
1050             SM_ERR_PRINT((TEXT("[SM :ERR] _ConstructSam fail!! (Vol %d, Part
1051             pstSMC->nVol, pstSMC->nPartID));
1052             SM_LOG_PRINT((TEXT("[SM :OUT] --SM_ReadSectors()\r\n\r\n")));
1053             return STL_CRITICAL_ERROR;
1054         }
1055
1056         while (nReadScts > 0)
1057         {
1058             pstCurrent = pstNew;
1059             nFirstOffset = 0xFFFFFFFF;
1060             nScts = 1;
1061             nReadScts--;
1062
1063             do
1064             {
1065                 if (pstCurrent->pSam[nSamIdx] < SM_SAM_DELETED)
1066                 {
1067                     /* get first sector offset */
1068                     nFirstOffset = pstCurrent->pSam[nSamIdx];
1069                     nSamIdx++;
1070
1071                     /* get the number of sequential sectors */
1072                     while (nReadScts > 0)
1073                     {
1074                         if ((nFirstOffset + nScts) == pstCurrent->pSam[nSamI
1075                         {
1076                             nScts++;
1077                             nReadScts--;
1078                             nSamIdx++;
1079                         }
1080                         else
1081                         {
1082                             break;
1083                         }
1084                     }
1085
1086                     /* read multiple sectors through BML */
1087                     nBErr = BML_MRead(pstVNC->nVol,
1088                                     pstSMC->pstSHPC->nStartVsn + nFirstOff
```

- MSR reads consecutive physical sectors together for improving read performance
- Statistics
 - 157 lines long
 - 4 level nested loops
 - 4 parameters to specify logical data to read (from where, to where, how long, read flag)

Loop Structure of MSR

```
01: curLU = LU0;
02: while( curLU != NULL ) { Loop1: iterates over LUs
03:     readScts = # of sectors to read in the current LU
04:     while( readScts > 0 ) { Loop2: iterates until the current LU is read completely
05:         curPU = LU->firstPU;
06:         while( curPU != NULL ) { Loop3: iterates over PUs linked to the current LU
07:             while(...) { Loop4: identify consecutive PS's in the current PU
08:                 conScts = # of consecutive PS's to read in curPU
09:                 offset = the starting offset of these consecutive PS's in curPU
10:             }
11:             BML_READ( curPU, offset, conScts );
12:             readScts = readScts - conScts;
13:             curPU = curPU->next;
14:         }
15:     }
16:     curLU = curLU->next;
17: }
```

PART II: Model Checking Results

- **Verification of MSR by using NuSMV, Spin, and CBMC**
 - NuSMV: BDD-based symbolic model checker
 - Spin: Explicit model checker
 - CBMC: C-bounded model checker
- **The requirement property is to check**
 - $\text{after_MSR} \rightarrow (\forall i. \text{logical_sectors}[i] == \text{buf}[i])$
- **We compared these three model checkers empirically**

Verification by NuSMV

- **NuSMV was the first choice as a verification tool, since**
 1. BDD-based symbolic model checkers have been known to handle large state spaces
 2. MSR operates with a semi-random environment (i.e. all possible configurations of PUs and SAMs analyzed)
 3. Data structure of MSR can be abstracted in a simple array form with assignments and equality checking operations only
 4. MSR is a single-threaded program

Target Model Creation in NuSMV

- We had to introduce control points variables, since
 - C is **control-flow based**
 - NuSMV modeling language is **dataflow-based**
- **Linked list is replaced by an array operation.**
 - Array index variables should be statically expanded, since NuSMV does not support index variables
- **As a result, the final NuSMV model is more than 1000 lines long**

A fragment of C	Conversion to parallel statements based on control and data dependency	Corresponding NuSMV code
<pre> 1: x=x-1; ← DP1 2: while(x>=0){ 3: y = x; ← DP2 4: x --;} ← DP3 </pre>	<pre> 0: DP1=0; DP2=0; DP3=0; 1: if (!DP1) { x=x-1; DP1 = 1;} 2: if ((DP1 DP3) && x>=0) { y = x; DP2=1; DP3=0; } 3: if (DP2) { x--; DP3=1; DP2=0; } </pre>	<pre> init(DP1):=0; init(DP2):=0; init(DP3):=0; next(DP1):= 1; next(DP2):= case (DP1 DP3) & (x >= 0) : 1; DP2 : 0; 1 : DP2; esac; next(DP3):= case (DP1 DP3) & (x >= 0) : 0; DP2 : 1; 1 : DP3; esac; next(x):= case !DP1 DP2 : x-1; 1 : x; esac; next(y):= case (DP1 DP3) & (x >= 0) : x; 1 : y; esac; </pre>

Modeling in NuSMV (2/2)

Environment model creation

- The environment of MSR (i.e., PUs and SAMs configurations) can be described by **invariant rules**. Some of them are

1. One PU is mapped to at most one LU
2. *Valid correspondence between SAMs and PUs:*

If the i th LS is written in the k th sector of the j th PU, then the i th offset of the j th SAM is valid and indicates the k 'th PS ,

Ex> 3rd LS ('C') is in the 3rd sector of the 2nd PU, then SAM1[2] ==2

$i=3$

$k=3$


$j=2$

3. *For one LS, there exists only one PS that contains the value of the LS:*


The PS number of the i th LS must be written in only one of the $(i \bmod 4)$ th offsets of the SAM tables for the PUs mapped to the corresponding LU.

	SAM0~SAM4				PU0~PU4			
Sector 0	1		0				E	
Sector 1		1		1	A	B		F
Sector 2		2				C		
Sector 3			3				D	

Verification Performance of NuSMV



(a) Time consumption



(b) Memory consumption

- Verification was performed on the machine equipped with Xeon5160 (3Ghz, 32Gbyte Memory), 64 bit Fedora Linux 7, NuSMV 2.4.3
- The requirement property was proved correct for all the experiments (i.e., MSR is correct in this small model)
 - For 7 sectors long data that are distributed over 7 PUs consumes more than 11 hours while consuming only 550 mb memory

Performance Analysis

- The MSR model (5 LS's and 5 PUs) has 365 BDD variables for its symbolic representation
 - At least 240 BDD variables are required for PUs and SAMs
 - $5 \text{ (# of PUs)} \times 4 \text{ (sectors/PU)} \times 2 \text{ (current/next)} \times 3 \text{ (bits)}$
- The same MSR model generated **1.2 million** BDD nodes.
- **Dynamic reordering** takes more than **90%** of total verification time
 - Time is the bottleneck in this NuSMV verification task

Modeling by Spin

- **A target model**


- Translated from the MSR C code through Modex which is an automated C-to-Promela translator with embedded C statements
 - Modex translates MSR into the same 4 level-nested loop control structure

- **An environment model**


- PUs and SAMs, which takes most of memory, are tracked, but **not** stored in the state vector through a **data abstraction technique**
 - `c_track` keyword and `Unmatched` parameter
 - Based on the observation that SAMs and PUs are sparse
 - Only a unique **signature** of the current state of PUs and SAMs is stored succinctly
 - $\langle (0,1), (1,1), (1,2), (2,3), (3,0), (4,1) \rangle$ is the signature of the following PUs and SAMs configuration

	SAM0~SAM4				PU0~PU4			
<i>Sector 0</i>	1			0			E	
<i>Sector 1</i>		1		1	A	B		F
<i>Sector 2</i>		2				C		
<i>Sector 3</i>			3				D	

Verification Performance of Spin



(a) Time consumption



(b) Memory consumption

- The requirement property was satisfied
- The data abstraction technique shows significant performance improvement upto **78%** of memory reduction and **35%** time reduction (for 5 logical sectors data)

# of physical units	5	6	7	8	9	10
Memory reduction	17%	38%	57%	68%	74%	78%
Time reduction	23%	24%	26%	32%	34%	35%


Modeling by CBMC

- CBMC does not require an explicit target model creation
- An environment for MSR was specified using **assume statements** and the environment model was similar to the environment model in NuSMV
- For the **loop bounds**, we can get valid upper bounds from the loop structure and the environment setting
 - The outermost loop: L times (L is a # of LUs)
 - The 2nd outermost loop: 4 times (one LU contains 4 LS's)
 - The 3rd outermost loop: M times (M is a # of PUs)
 - The innermost loop: 4 times (one PU contains 4 PS's)


L=2, M=5

	SAM0~SAM4				PU0~PU4			
Sector 0	1		0				E	
Sector 1		1		1	A	B		F
Sector 2		2				C		
Sector 3		3					D	

Verification Performance of CBMC





(a) Time consumption



(b) Memory consumption

- Exponential increase in both time and memory. However, the slope is much lower than those of NuSMV and Spin, which makes CBMC perform better for large problems
- A problem of 10 PUs and 8 LS's has 8.6×10^5 variables and 2.9×10^6 clauses.

Performance Comparison



Conclusion

- **Application of Model Checking to Industrial SW Project**
 - Current off-the-shelf model checkers showed their effectiveness to debug a part of industrial software, if a target portion is carefully selected
 - Although model checker worked on a small scale problem, it still contributes due to its exhaustive exploration which is complementary to the testing result
- **Comparison among the Three Model Checkers**

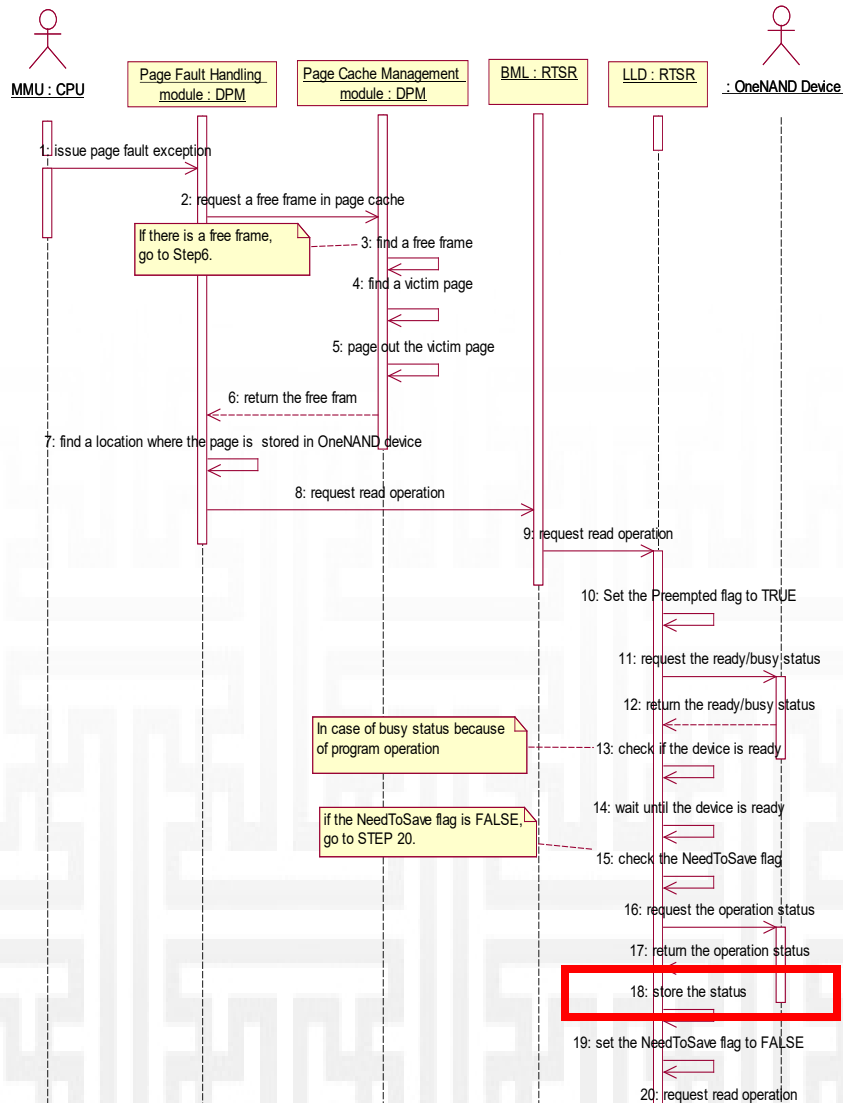
	Modeling Difficulty	Memory Usage	Verification Speed
NuSMV	Most difficult	Good	Slow
Spin	Medium difficult	Poor	Fast
CBMC	Easiest	Best	Fastest



Requirement Elicitation

- Utilizes a top-down approach
 - We selected 3 functional requirements with “Very High” priority in the USP SRS document
 - “2.2.1 Support prioritized read operation” (pg 14)
 - Issues on **synchronization**
 - “2.2.3 Keep data integrity” (pg 15)
 - Issues on **handling exceptions/failures**
 - “2.2.6 Manage sectors” (pg 17)
 - Issues on **complex sector mapping**
 - Focused on the above three requirements, we investigated STL/BML/LLD/DPM documents as well as C code to enumerate detailed points to check explicitly
 - Reported in the 50 pgs document “**Requirement Specification과 Detailed Design Specification 분석을 통한 검증기준 정리 결과**”

Ex> Page Fault Handling While a Device is Being Programmed (ADS 4.4 p.22)



....

13. LLD checks whether OneNAND is ready status or busy status.

- 현재 busy 상태임을 정확히 인식하는지 확인

14. If OneNAND is busy status, wait until it is ready.

- 무한 loop에 빠지는 상황발생 가능성 확인

15. After OneNAND become ready status, it checks whether the NeedToSave Flag is set to TRUE or FALSE.

16. If the NeedToSave Flag is set to TRUE, LLD checks a program or erase operation status of OneNAND.

17. OneNAND returns the program status or the erase status.

- CTRL_STAT 상태값의 정확한 반환확인

18. LLD stores the program status or the erase status.

- CTRL_STAT 상태값을 nSavedStatus에 저장여부확인

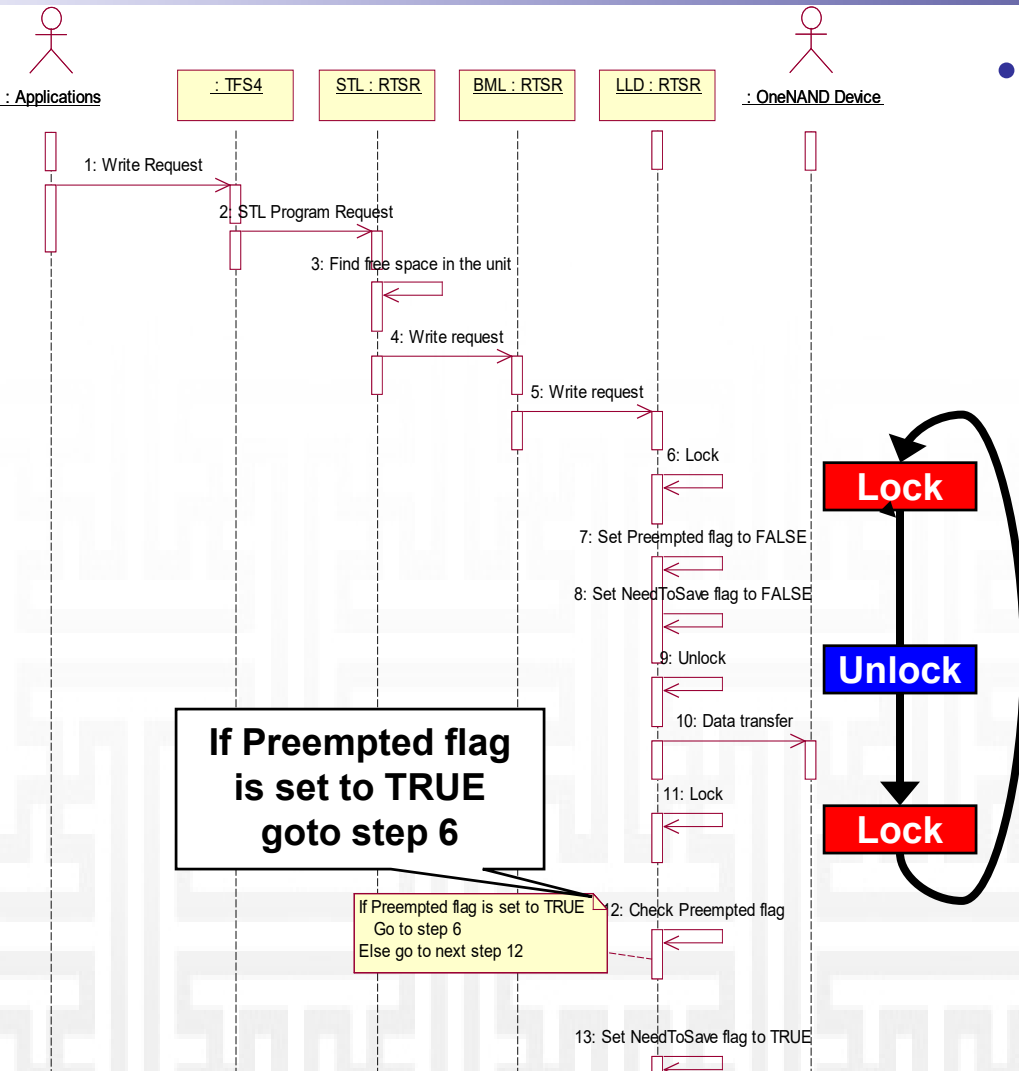
19. LLD sets the NeedToSave Flag to FALSE

20. LLD requests OneNAND to read the desired page.

21~23. Return Success/Fail.

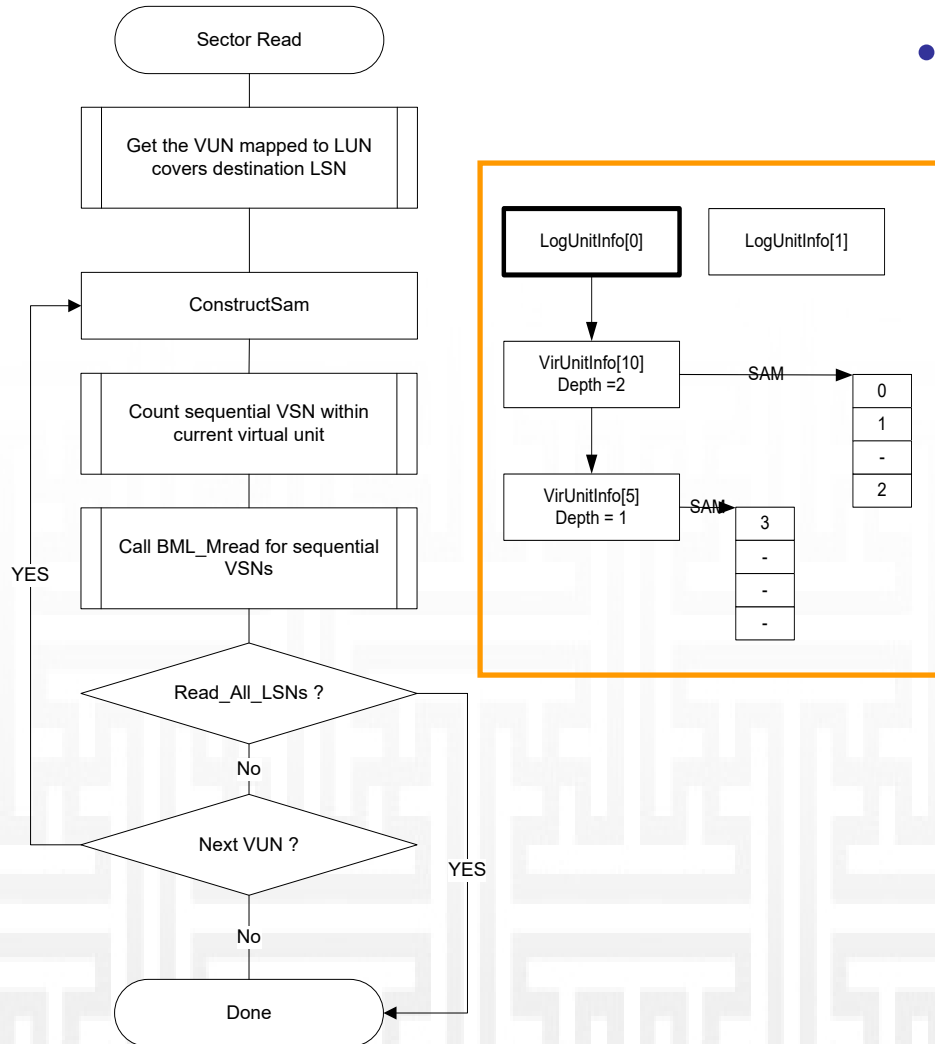
...

Ex 1. Inconsistencies Detected



- “Data writing without page fault” ADS 4.8 28pg
 - It describes a use case where repeated locking without unlocking is possible
 - However, actual code (`ONLD_Write()` in `ONLDNonAtomic.c`) does not have such behavior because unlock is performed before jumping back

Ex 2.Inconsistencies Detected



- “Sector Read” STL DDS 4.2.4 34pg

- It describes that multisector reading operation traverses a list of PU only once, which causes incomplete sector read
- However, actual code (`SM_ReadSectors()` in `SectorMap.c`) traverses a list of PU repeatedly, which enables complete sector reading with a performance penalty

Figure 4-11 The flow chart to read sectors
Formal Verification of a Flash Memory Device
Driver – an Experience Report

Excerpts of the SMV Model

MODULE main

-- Variable declaration

VAR

SAM : array 0..4 of sam_type;

PU : array 0..4 of PU_type;

buf : array 0..4 of 0..5;

nScts : 0..5;

-- SPEC

INVARSPEC (after_first_do ->

PU[0].sect[0]=1 &

PU[0].sect[1]=2 &

PU[0].sect[2]=3 &

PU[0].sect[3]=4 &

PU[3].sect[0]=5)

init(buf[0]):=0;

-- if(pBuf==0 && 0 < nScts)

-- buf[0]= PU[PU_id].sect[nFirstOffset]

next(buf[0]):

case after_fourth_do :

case pBuf = 0 & 0 < nScts: -- i=0

case

PU_id=0 & nFirstOffset=0: PU[0].sect[0];

PU_id=0 & nFirstOffset=1: PU[0].sect[1];

PU_id=0 & nFirstOffset=2: PU[0].sect[2];

PU_id=0 & nFirstOffset=3: PU[0].sect[3];

...

PU_id=4 & nFirstOffset=3 : PU[4].sect[3];

esac;

esac;

init(buf[1]):=0;

next(buf[1]):= ...

Excerpts of the Spin Model

```
active proctype SM_ReadSectors() {
```

```
    byte buf[NUM_LS_USED];
    byte nScts;
    byte nFirstOffset;
    byte nNumOfScts=NUM_LS_USED;
    byte nReadScts=nNumOfScts;
    byte nSamIdx;
```

```
    do /* 1047: while (nNumOfScts >0) { */
```

```
        :: nNumOfScts > 0 ->
            PU_id = lui[nLun];
            if /* nReadScts = ... */
                :: (SECT_PER_U-nSamIdx)> nNumOfScts ->
                    nReadScts = nNumOfScts;
                :: else->nReadScts =SECT_PER_U- nSamIdx;
            fi;
            nNumOfScts = nNumOfScts - nReadScts;
```

```
        do /* line 1068: while (nReadScts > 0) */
            :: (nReadScts > 0) -> PU_id = lui[nLun];
                nFirstOffset=255;
                nScts=1; nReadScts--;
```

```
        do /* line 1075: do {... */
```

```
            :: true;
            if /* line 1077: if(pstCurrent->pSam[nSamIdx]...*/
                :: SAM[PU_id].valid[nSamIdx]-> nFirstOffset =
                    SAM[PU_id].offset[nSamIdx];nSamIdx++;
                do /* line 1084:while (nReadScts > 0) { ...} */
                    :: (nReadScts > 0) ->
                        if
                            ::FirstOffset+nScts==
                                SAM[PU_id].offset[nSamIdx] ->
                                    nScts++;nReadScts--;nSamIdx++;
                                :: else-> break;
                        fi;
                    :: else->break;
                od;
```

```
                BML_MRead(PU_id,nFirstOffset,nScts,pBuf);
                break;
```

```
            :: else;
            fi;
            if /*line 1112: } while ( PU[PU_id].nil != true) */
                :: PU[PU_id].nil -> break;
            :: else;
            fi;
            PU_id++;
```

```
        od;
    }
```

Difficulties in Testing the USP File System

- **Limitation of detecting bugs through actual testing**


- For example, if data is l sectors long and there are n physical units, then the possible number of distribution is

$$\sum_{i=1}^{n-1} ((4 \times i) C_4 \times 4!) \times ((4 \times (n-i)) C_{(l-4)} \times (l-4)!)$$

- If 6 LS's are distributed over 1000 PUs, 3.9×10^{22} cases exist
 - Even 10^{10} random tests only covers tiny fraction of possible cases

- **Lack of fine control over the USP system**

- Most testings are performed at the file system level
 - Subtle bugs are hard to detect, for example



Modeling in NuSMV (2/2)

- **Environment model creation**

- The environment of MSR (i.e., PUs and SAMs configurations) can be described by invariant rules. Some of them are
 1. One PU is mapped to at most one LU
 2. If the i th LS is written in the k th sector of the j th PU, then the $(i \bmod 4)$ th offset of the j th SAM is valid and indicates the PS number k ,
- 1. The PS number of the i th LS must be written in only one of the $(i \bmod 4)$ th offsets of the SAM tables for the PUs mapped to the corresponding LU.

$$\begin{aligned} \forall i, j, k \ (logical_sectors[i] = PU[j].sect[k] \rightarrow (SAM[j].valid[i \bmod m] = true \\ \& SAM[j].offset[i \bmod m] = k \\ \& \forall p. (SAM[p].valid[i \bmod m] = false) \\ \text{where } p \neq j \text{ and } PU[p] \text{ is mapped to } \lfloor \frac{i}{m} \rfloor_{th} \text{ LU})) \end{aligned}$$

USP Code Statistics

