

Greybox Fuzzing

Shin Hong

CSEE Handong Global Univ.

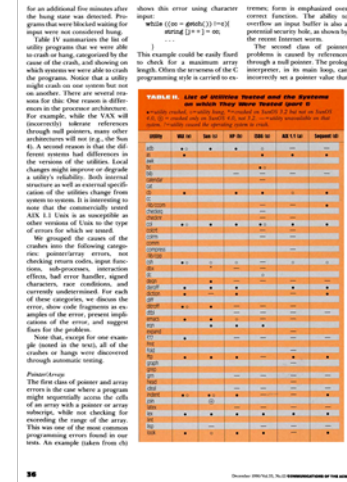

Topics

- fuzzing background
 - mutation-based fuzzing
 - greybox fuzzing
- introduction to the libFuzzer tool
 - functionalities
 - tool structure
 - walkthrough example
- engineering aspects of unit test fuzzing

It was a Dark and Stormy Night in the Fall of 1988

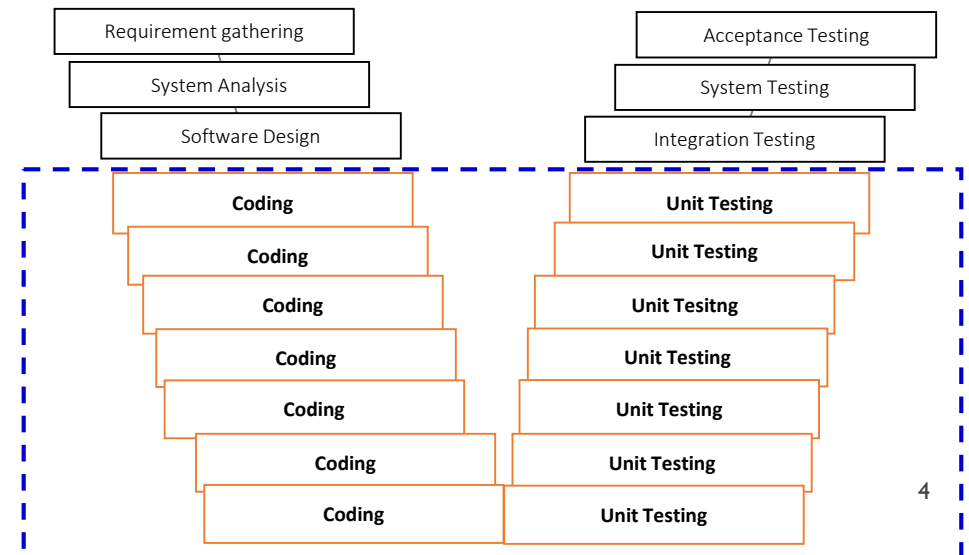
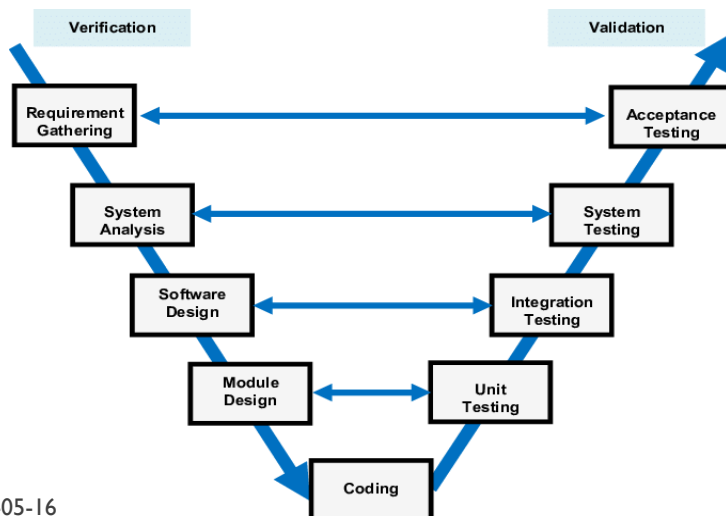
<http://pages.cs.wisc.edu/~bart/fuzz/ForewordI.html>

- Barton Miller, a professor of U. Wisconsin-Madison experienced that UNIX systems crashed extraordinary frequently.
- He conjectured that it was because unexpectedly strong electric noise induced multiple tweaks in packets
- To test his conjecture, Miller gave an assignment to students to test UNIX utilities by feeding intentionally randomized inputs
 - Miller et al., An empirical study of the reliability of UNIX utilities, CACM, 1990



Ancient Fuzzers

- Generate a long sequence of random texts that have similar aspects as formatted text input for testing UNIX command utilities
 - intermix comma, semicolon, and many control characters
 - e.g., `'!7#%"*#0=)$;%6*;>638:*>80"=`
 - Feed randomly generated texts to a target UNIX utility, and repeat this for many hours
- By using this kind of ancient fuzzers, new bugs were found from **one third** of the UNIX utilities



Shortcomings of Ancient Fuzzers

- Ancient fuzzers detect only crashes and hangs, but cannot uncover **silent illegal behaviors** which can result much critical consequences
 - reliability issue \Rightarrow security issue (adversarial users)
 - employ dynamic analyzers to detect and/or predict silent violations
 - e.g., valgrind, electric fence, LLVM sanitizer suites (AddressSanitizer, MemorySanitizer, UndefinedBehaviorSanitizer)
- Randomly generated inputs **cover only restricted portion of the source code**
 - random inputs are often rejected quickly because they likely have trivial input grammar errors
 - extremely low probability for a randomly generated text to pass grammar checks

Mutation-based Fuzzing

- Ideas

- start with a set of **valid inputs** (seeds)
- repeatedly introduce **small changes** to the existing inputs (*mutation*) with a hope that they exercise new behaviors

- Example: fuzzing a URL parsing library

- Seed

- `http://www.google.com/search?q=fuzzing`

- Fuzzed inputs

- `http://www.g=0Nogl.om/search?q=fuzzing/`
- `RttpX://w)ww.goo(gle.comq/sarc(q=fuzzng`
- `hdt8p:// "wWw.goole.com/seDarb`*?q=fuzzing`
- `hup://www.google.comC/search?q=fuzzing`
- `http://w7w.google.com/search?q=ufuzgzng`
- `http://w&ww.google.cKom/search7q=fuzzing`

Mutation Operators

- Flip one random bit
- Alternate one or multiple consecutive bytes
- Erase one or multiple bytes from random offsets
- Insert one or multiple bytes to random offsets
- Repeat existing bytes multiple times
- Add a word from a predefined dictionary
- Shuffle consecutive bytes (reorder multiple bytes randomly)
- Copy a substring and paste it randomly offsets
- Crossover
- Apply mutation one or more times on a single seed input

Fine-grained



Coarse-grained

Why Mutation Effectively Disclose Subtle Behaviors?

- It is likely to obtain quality seed inputs from existing test cases
- An error-revealing input mostly resides close to a valid input
 - close in lexical distance, or numerical distance
 - competent programmer hypothesis
- A part of a program input is likely associated with only few program components
 - an aspect of an input text can be represented as a short subsequence
 - strong locality exists in a well-modularized program
- A critical value of a specific part of input is likely found in the other parts of the inputs

Greybox Fuzzing: Use Structural Coverage to Guide Fuzzing

- Idea

- Start with a set of valid inputs
- Repeatedly introduce small changes to the existing inputs while expecting they exercise new behaviors
- Include the mutated input as a seed only if it explores a new behavior
 - covering a new structural test requirement

- Greybox fuzzers (e.g., AFL, libFuzzer) show in practice that use of structural coverage dramatically improves effectiveness of mutation-based fuzzing
 - Google runs fuzzing on 160 open-source projects with 250,000 machines
 - Google found more than 16,000 bugs in Chrome by fuzzing

Basic Algorithm

Input: a target program $Prog$
a set of seeds $S = \{s_1, s_2, \dots, s_n\}$
Output: two sets of tests $P = \{p_1, p_2, \dots, p_m\}$, $F = \{f_1, f_2, \dots, f_k\}$

Procedure:

```
 $P \leftarrow S, F \leftarrow \emptyset, C \leftarrow \emptyset$ 
while  $p \in P$  begin
     $C \leftarrow C \cup \text{Cov}(Prog, p)$ 
end while
while termination condition is not satisfied begin
     $p \leftarrow$  select a random test input from  $P$ 
     $p' \leftarrow$  mutate  $p$  with a certain mutation operator
    if  $Prog(p')$  fails then
         $F \leftarrow F \cup \{p'\}$ 
    else
        if  $\text{Cov}(Prog, p') - C \neq \emptyset$  then
             $P \leftarrow P \cup \{p'\}$ 
             $C \leftarrow C \cup \text{Cov}(Prog, p')$ 
        end if
    end if
end while
```

libFuzzer: Fuzzing Tool for LLVM

<https://llvm.org/docs/LibFuzzer.html>

- libFuzzer is a greybox fuzzer inspired by AFL for testing C/C++ libraries
 - developed as a component of LLVM
 - target C/C++ programs
 - well integrated with the LLVM sanitizer suites
 - generate inputs to public APIs in a unit test driver (rather than a system input)
 - provide a plugin API for defining and managing **custom mutation operators**
 - easy to implement structure-aware, grammar-based fuzzing
- libFuzzer, together with AFL, is used as a core component of OSS-Fuzz and ClusterFuzz <https://google.github.io/clusterfuzz/>



libFuzzer Mutation Operators

Mutator	Description
EraseBytes	Reduce size by removing a random byte
InsertByte	Increase size by one random byte
InsertRepeatedBytes	Increase size by adding at least 3 random bytes
ChangeBit	Flip a Random bit
ChangeByte	Replace byte with random one
ShuffleBytes	Randomly rearrange input bytes
ChangeASCII Integer	Find ASCII integer in data, perform random math ops and overwrite into input.
ChangeBinary Integer	Find Binary integer in data, perform random math ops and overwrite into input
CopyPart	Return part of the input
CrossOver	Recombine with random part of corpus/self
AddWordPersist AutoDict	Replace part of input with one that previously increased coverage (entire run)
AddWordTemp AutoDict	Replace part of the input with one that recently increased coverage
AddWord FromTORC	Replace part of input with a recently performed comparison

- Domain-specific word dictionary can be configured for a specific target function
- We can add custom mutation operators
 - alternate an input text considering its grammar or constraints on input validity

Writing Unit Fuzzing Driver (parameterized unit test case)

- *target function* accepts array of bytes, and feed accepted data into the API under test

```
// target.cc
extern "C" int LLVMFuzzerTestOneInput(const uint8_t *Data, size_t Size) {
    DoSomethingInterestingWithMyAPI(Data, Size);
    return 0; // Non-zero return values are reserved for future use.
}
```

- aspects
 - set prerequisite environment to run target API
 - configure test execution environment
 - invoke other APIs to set the starting state and also mock objects
 - cast given fuzzed input to the arguments of a target API
 - typecasting (e.g., a region of string to an integer)
 - precondition checking
 - selecting sub-cases of a test scenario
 - configure fuzzing engine

Example - Triangle

- Fuzzing target

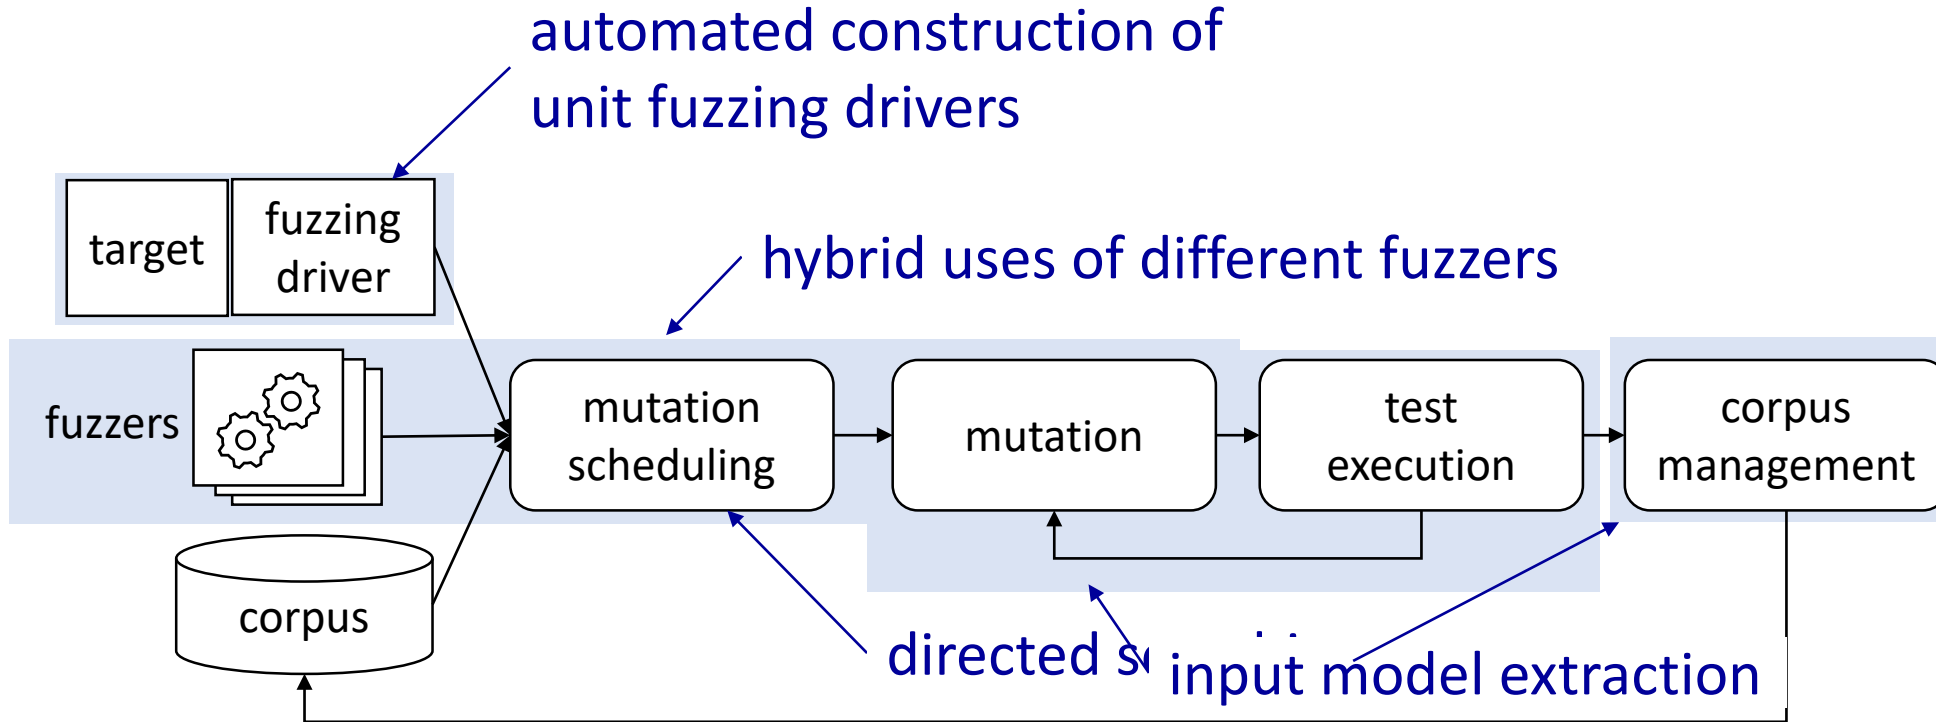
```
extern int LLVMFuzzerTestOneInput (const uint8_t *Data, size_t Size) {  
    if (Size != 12)  
        return 0 ;  
    int a, b, c ;  
    a = *((int *) (Data + 0)) ; b = *((int *) (Data + 4)) ; c = *((int *) (Data + 8)) ;  
  
    triangle_type(a, b, c) ;  
}
```

- Commands

```
clang -g -O1 -fsanitize=fuzzer,signed-integer-overflow triangle.c fuzz_target.c -o test-target  
export UBSAN_OPTIONS=halt_on_error=1  
./test-target corpus seed -max_len=100
```

Advanced Fuzzing Techniques

Recent Advances of Fuzz Testing Techniques



Unit Fuzzing Driver Generation

- Challenges: unit fuzzing requires **fuzzing drivers** for testing an API function, but writing a unit fuzzing driver is laborious and tricky
 - soundness issues
 - comprehensiveness issue
- Approaches
 - **FUDGE** extracts a fuzzing driver of a specific function via slicing from where the function is used in the production code [ESEC/FSE'19]
 - **FuzzGen** infers API dependency graphs from the source code and synthesizes likely API sequences as fuzzing drivers [USENIX Sec'20]
 - **G-EvoSuite** generates unit fuzzing drivers by extending their method sequences according to coverage feedback [ASE'20]

Hybrid Use of Different Fuzzers

- Challenges: no single golden fuzzer exists – the effectiveness of a certain fuzzer varies depending on target program characteristics
- Approaches: intermix different fuzzers and test generation techniques
 - **Zest** cooperates grammar-aware input construction and mutation fuzzing to induce synergetic effects [ISSTA'19]
 - **EnFuzzer** collaboratively runs multiple fuzzers via seed synchronization [USENIX-Sec'19]
 - **MUEZZ** learns an effective seed scheduling scheme for hybrid uses of greybox fuzzing and concolic testing techniques [RAID'20]
 - **HFL** combines fuzzing and symbolic execution for testing kernel code [NDSS'20]
 - **CUPID** automatically suggests for a combination of different fuzzers, that provide reliable performance to various targets [ACSAC'20]

Directed Searching

- Challenge: conventional fuzzers simply explores diverse paths without considering their values under testing context
- Approaches
 - **ParmaSan** gives high priorities to the seeds that cover more sanitizer warnings for detecting memory errors quickly [USENIX-Sec'20]
 - **IJSON** receives and interprets developer's annotations in guiding a fuzzing process [SP'20]
 - **Wustholz et al.** utilizes static analyses to predict likely paths to target code locations and uses the prediction as fuzzing guidance [ICSE'20]
 - **TortoiseFuzzer** assigns higher weights to the code that seem relevant to memory corruption faults to avoid searching meaningless paths [NDSS'22]
 - **TOFU** prioritizes seeds that reaches to closer paths to the given target code locations [preprint]

Input Model Extraction

- Challenges: fuzzers must be aware of the constraints on structured inputs, but it is hard for users to specify input grammar correctly and completely
- Approaches
 - **LFuzzer** extracts useful input tokens online and uses the extracted tokens within the fuzzing process [ISSTA'20]
 - **FuzzGuard** trains a CNN to predict if an input will reach to a target location and uses the model to reject useless inputs [USENIX-Sec'20]
 - **MTFuzz** trains a neural model to predict which coverage items are covered from input offsets and uses the model to guide mutations [ESEC/FSE'20]

References

The Fuzzing Book: Tools and Techniques for Generating Software Tests

Andreas Zeller, Rahul Gopinath, Marcel Böhme, Gordon Fraser, and Christian Holle

<https://www.fuzzingbook.org/>

libFuzzer – a library for coverage-guided fuzz testing

<https://llvm.org/docs/LibFuzzer.html>

The Art, Science and Engineering of Fuzzing: A Survey

V. J.M. Manes, H. Han, C. Han, S. K. Cha, M. Egele, E. J. Schwartz, and M. Woo

IEEE Transactions on Software Engineering, Volume: 47, Issue: 11, 01 November 2021